

GitHub Copilot Prompt: EV Battery Digital Twin System

Project Overview

Build a production-grade Electric Vehicle Battery Digital Twin system with:

- Real-time telemetry simulation
- AI-powered RUL (Remaining Useful Life) prediction using XGBoost
- Time-series data storage with TimescaleDB
- Message streaming via Kafka (Redpanda) and MQTT
- Live monitoring with Prometheus and Grafana
- MLflow for model tracking
- 3D battery visualization

Architecture Requirements

Tech Stack

- **Language:** Python 3.12+
- **ML Framework:** XGBoost, scikit-learn, NumPy, Pandas
- **Database:** TimescaleDB (PostgreSQL 15)
- **Message Brokers:** Redpanda (Kafka), Eclipse Mosquitto (MQTT)
- **Monitoring:** Prometheus, Grafana
- **ML Ops:** MLflow, MinIO (S3-compatible)
- **Containers:** Docker Compose

Project Structure

```
ev-battery-digital-twin/
├── src/
│   ├── __init__.py
│   ├── simulator/
│   │   ├── __init__.py
│   │   └── publisher.py      # Telemetry generator
│   ├── inference/
│   │   ├── __init__.py
│   │   └── live_predictor.py  # Real-time ML predictions
│   └── models/
│       ├── __init__.py
```

```
|   |   └── train.py      # XGBoost training pipeline
|   └── common/
|       ├── __init__.py
|       └── utils.py      # Shared utilities
|   └── models/          # Trained model artifacts
|   └── datasets/        # Training data (Kaggle dataset)
|   └── notebooks/       # Jupyter notebooks
|   └── infra/
|       ├── grafana/provisioning/
|       |   ├── datasources/
|       |   |   └── datasource.yml
|       |   └── dashboards/
|       |       ├── dashboard.yml
|       |       └── battery_dashboard.json
|       └── prometheus/
|           └── prometheus.yml
|   └── mosquitto/
|       └── mosquitto.conf
|   └── create_telemetry_table.sql
|   └── tests/            # Unit tests
|       └── docker-compose.yml
|   └── requirements.txt
|   └── .env.example
└── README.md
```

Implementation Instructions

1. Docker Compose Infrastructure

Create `docker-compose.yml` with these services:

TimescaleDB (port 5432):

- Image: `timescale/timescaledb:latest-pg15`
- Environment: `POSTGRES_USER=twin`, `POSTGRES_PASSWORD=twin_pass`, `POSTGRES_DB=twin_data`
- Volume: Init script at `/docker-entrypoint-initdb.d/init.sql`

Redpanda (ports 9092, 29092):

- Image: `docker.redpanda.com/redpandadata/redpanda:v24.2.4`
- Kafka-compatible message broker
- Topic: `ev-telemetry`

MinIO (ports 9000, 9001):

- Image: `minio/minio:latest`
- S3-compatible storage for ML artifacts

MLflow (port 5000):

- Image: `ghcr.io/mlflow/mlflow:v2.9.2`
- Backend: SQLite
- Artifact store: MinIO S3

Mosquitto (ports 1883, 9002):

- Image: `eclipse-mosquitto:2.0`
- MQTT broker for IoT devices

Prometheus (port 9090):

- Image: `prom/prometheus:latest`
- Scrapes metrics from predictor service

Grafana (port 3000):

- Image: `grafana/grafana:latest`
- Default credentials: admin/admin
- Auto-provision TimescaleDB datasource

Network: `twin-net` (bridge)

2. Database Schema (`infra/create_telemetry_table.sql`)

Create TimescaleDB hypertable:

```
sql
```

```
CREATE EXTENSION IF NOT EXISTS timescaledb;
```

```
CREATE TABLE IF NOT EXISTS ev_telemetry (
    time TIMESTAMPTZ NOT NULL,
    battery_id VARCHAR(50) NOT NULL,
    soc DOUBLE PRECISION,
    soh DOUBLE PRECISION,
    voltage DOUBLE PRECISION,
    current DOUBLE PRECISION,
    temperature DOUBLE PRECISION,
    charge_cycles INTEGER,
    power_consumption DOUBLE PRECISION,
    rul_prediction DOUBLE PRECISION,
    failure_probability DOUBLE PRECISION,
    PRIMARY KEY (time, battery_id)
);
```

```
SELECT create_hypertable('ev_telemetry', 'time', if_not_exists => TRUE);
CREATE INDEX idx_battery_id ON ev_telemetry(battery_id, time DESC);
SELECT add_retention_policy('ev_telemetry', INTERVAL '90 days', if_not_exists => TRUE);
```

3. ML Model Training ([\(src/models/train.py\)](#))

Implement XGBoost training pipeline:

Features (7 inputs):

1. SoC (State of Charge) %
2. SoH (State of Health) %
3. Battery_Voltage (V)
4. Battery_Current (A)
5. Battery_Temperature (°C)
6. Charge_Cycles (count)
7. Power_Consumption (W)

Target Variables:

- RUL (Remaining Useful Life) - Regression
- Failure_Probability - Binary Classification

Pipeline:

1. Load dataset from `(datasets/EV_Predictive_Maintenance_Dataset_15min.csv)`

2. Split data (80/20 train/test)

3. StandardScaler for feature normalization

4. Train XGBRegressor for RUL (target $R^2 > 0.90$)

5. Train XGBClassifier for failure probability

6. Log experiments to MLflow

7. Save models using joblib:

- `rul_xgb_model.joblib`

- `failure_xgb_model.joblib`

- `scaler.joblib`

- `feature_names.joblib`

Hyperparameters:

- `n_estimators: 200`

- `max_depth: 8 (RUL), 6 (Failure)`

- `learning_rate: 0.1`

- `subsample: 0.8`

- `colsample_bytree: 0.8`

4. Telemetry Simulator (`(src/simulator/publisher.py)`)

Implement realistic EV battery simulation:

Behavior:

- Discharge: SoC decreases 0.2-0.8% per cycle

- Charge: SoC increases 0.5-1.5% per cycle

- Switch to charging at SoC < 20%

- Complete charge at SoC > 99%

- SoH degrades 0.01% per cycle

- Temperature rises with current draw

Publishing:

1. PostgreSQL/TimescaleDB - Store telemetry

2. Kafka topic `(ev-telemetry)` - Stream processing
3. MQTT topic `(ev/battery/telemetry)` - IoT devices

Interval: Every 2 seconds (configurable)

Output format (JSON):

```
json

{
  "timestamp": "2025-10-20T10:30:15.123Z",
  "battery_id": "BATTERY_001",
  "soc": 85.5,
  "soh": 95.2,
  "voltage": 380.5,
  "current": 105.3,
  "temperature": 32.1,
  "charge_cycles": 150,
  "power_consumption": 40.1
}
```

5. Live Prediction Service (`(src/inference/live_predictor.py)`)

Implement real-time ML inference:

Workflow:

1. Load trained models from `(models/)` directory
2. Every 5 seconds:
 - Fetch latest telemetry from TimescaleDB
 - Scale features using saved scaler
 - Predict RUL and failure probability
 - Update database with predictions
 - Expose Prometheus metrics

Prometheus Metrics (port 9100):

- `(battery_rul_prediction)` (Gauge)
- `(battery_failure_probability)` (Gauge)
- `(battery_soc)` (Gauge)
- `(battery_soh)` (Gauge)

- `battery_temperature` (Gauge)
- `predictions_total` (Counter)

Dependencies:

- joblib for model loading
- psycopg2 for database
- prometheus_client for metrics

6. Prometheus Configuration ([\(infra/prometheus/prometheus.yml\)](#))

```
yaml

global:
  scrape_interval: 5s

scrape_configs:
  - job_name: 'ev_predictor'
    static_configs:
      - targets: ['host.docker.internal:9100']
```

7. Grafana Dashboard Configuration

Create datasource ([\(infra/grafana/provisioning/datasources/datasource.yml\)](#)):

```
yaml

apiVersion: 1
datasources:
  - name: TimescaleDB
    type: postgres
    url: timescaledb:5432
    database: twin_data
    user: twin
    secureJsonData:
      password: twin_pass
    jsonData:
      sslmode: disable
      timescaledb: true
```

Create dashboard with panels:

1. SoC Time Series - Line chart

2. SoH Trend - Line chart with threshold

3. Temperature Heatmap - Color-coded

4. RUL Gauge - Current prediction

5. Failure Risk - Percentage gauge

6. Charge Cycles Counter - Stat panel

7. Voltage/Current Graph - Dual-axis

8. Power Consumption - Area chart

Auto-refresh: 5 seconds

8. Dependencies ([\(requirements.txt\)](#))

```
# Core ML & Data Science
```

```
xgboost>=2.0.0
```

```
numpy>=1.24.0
```

```
pandas>=2.0.0
```

```
scikit-learn>=1.3.0
```

```
joblib>=1.3.0
```

```
# ML Ops
```

```
mlflow>=2.9.0
```

```
boto3>=1.28.0
```

```
# Database
```

```
psycopg2-binary>=2.9.0
```

```
# Messaging
```

```
kafka-python>=2.0.0
```

```
paho-mqtt>=1.6.0
```

```
# Monitoring
```

```
prometheus-client>=0.18.0
```

```
# Web Framework
```

```
flask>=3.0.0
```

```
flask-cors>=4.0.0
```

```
# Utilities
```

```
python-dotenv>=1.0.0
```

```
pyyaml>=6.0
```

```
requests>=2.31.0
```

```
# Development
jupyter>=1.0.0
matplotlib>=3.7.0
seaborn>=0.12.0
```

9. Environment Configuration ([\(.env.example\)](#))

```
bash

# Database
PG_HOST=localhost
PG_PORT=5432
PG_USER=twin
PG_PASSWORD=twin_pass
PG_DATABASE=twin_data

# Kafka
KAFKA_BOOTSTRAP_SERVERS=localhost:9092
KAFKA_TOPIC=ev-telemetry

# MQTT
MQTT_HOST=localhost
MQTT_PORT=1883
MQTT_TOPIC=ev/battery/telemetry

# MLflow
MLFLOW_TRACKING_URI=http://localhost:5000
MLFLOW_S3_ENDPOINT_URL=http://localhost:9000
AWS_ACCESS_KEY_ID=minioadmin
AWS_SECRET_ACCESS_KEY=minioadmin

# Model paths
MODEL_DIR=models
```

10. Utility Functions ([src/common/utils.py](#))

Implement helper functions:

```
python
```

```

def get_db_connection():
    """Create PostgreSQL connection with retry logic"""
    pass

def create_kafka_producer():
    """Initialize Kafka producer with error handling"""
    pass

def create_mqtt_client():
    """Setup MQTT client with connection callbacks"""
    pass

def load_models(model_dir):
    """Load all trained models and scaler"""
    pass

def validate_telemetry(data):
    """Validate telemetry data structure"""
    pass

def setup_logging(name, log_file='logs/app.log'):
    """Configure structured logging"""
    pass

```

11. Testing ([tests/test_simulator.py](#), [tests/test_predictor.py](#))

Create unit tests:

- Test telemetry generation logic
- Test database insertion
- Test model prediction accuracy
- Test metric exposition
- Mock external services

12. Documentation ([README.md](#))

Include:

- Architecture diagram
- Quick start guide

- Service endpoints table
 - Model performance metrics
 - Troubleshooting section
 - API documentation
 - Development setup
-

Execution Flow

Setup & Deployment

```
bash
```

```
# 1. Clone and setup
```

```
git clone <repo>
cd ev-battery-digital-twin
python3.12 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt
```

```
# 2. Download dataset from Kaggle
```

```
# Place in: datasets/EV_Predictive_Maintenance_Dataset_15min.csv
```

```
# 3. Start infrastructure
```

```
docker-compose up -d
```

```
# 4. Wait for services (30 seconds)
```

```
sleep 30
```

```
# 5. Train models
```

```
python src/models/train.py
```

```
# 6. Start simulator (Terminal 1)
```

```
python src/simulator/publisher.py
```

```
# 7. Start predictor (Terminal 2)
```

```
python src/inference/live_predictor.py
```

```
# 8. Access dashboards
```

```
# Grafana: http://localhost:3000 (admin/admin)
```

```
# MLflow: http://localhost:5000
```

```
# Prometheus: http://localhost:9090
```

```
# MinIO: http://localhost:9001
```

Advanced Features to Implement

1. REST API ([src/api/server.py](#))

Flask application with endpoints:

- `GET /api/telemetry/latest` - Latest battery data
- `GET /api/telemetry/history?minutes=60` - Historical data
- `GET /api/predictions/current` - Current RUL and failure risk
- `GET /api/health` - Service health check
- `POST /api/battery/configure` - Update simulation parameters

2. WebSocket Real-time Stream

Live data streaming to web clients using Flask-SocketIO

3. Anomaly Detection

Implement Isolation Forest for detecting abnormal patterns in telemetry

4. Alert System

Trigger alerts when:

- SoC < 15% (critical low battery)
- Temperature > 60°C (overheating)
- Failure probability > 0.7 (high risk)
- RUL < 100 cycles (maintenance needed)

Send notifications via:

- Email (SMTP)
- Slack webhook
- Database log table

5. Multi-Battery Fleet Management

Extend to handle multiple batteries:

- Array of battery_ids
- Fleet-level dashboard
- Comparative analytics

6. Model Retraining Pipeline

Scheduled retraining:

- Collect new telemetry data
- Retrain models monthly
- Compare performance metrics
- Auto-deploy if improved

7. Data Export Service

Export telemetry to:

- CSV files
- Parquet format
- Cloud storage (S3, Azure Blob)

8. Load Testing

Simulate 100+ batteries with stress testing tools

Code Quality Requirements

Style Guide

- Follow PEP 8
- Type hints for all functions
- Docstrings (Google style)
- Maximum line length: 100 characters

Error Handling

- Try-except blocks for all I/O operations
- Custom exception classes
- Graceful degradation
- Detailed error logging

Logging

- Structured JSON logs
- Log levels: DEBUG, INFO, WARNING, ERROR, CRITICAL

- Separate log files per service
- Log rotation (max 10MB per file)

Security

- No hardcoded credentials (use environment variables)
 - SQL injection prevention (parameterized queries)
 - Input validation on all external data
 - Rate limiting on API endpoints
-

Performance Targets

- **Telemetry Generation:** 30+ records/minute
 - **Prediction Latency:** < 100ms per inference
 - **Database Write:** < 50ms per insert
 - **Model Loading:** < 3 seconds on startup
 - **Memory Usage:** < 500MB per service
 - **API Response Time:** < 200ms (p95)
-

Monitoring & Observability

Metrics to Track

1. System Metrics:

- CPU usage per service
- Memory consumption
- Disk I/O
- Network throughput

2. Application Metrics:

- Telemetry generation rate
- Prediction frequency
- Model inference time
- Database query duration
- Message queue lag

3. Business Metrics:

- Average battery SoH
- Predicted RUL distribution
- Failure risk trends
- Charge cycle statistics

Grafana Alerts

Configure alert rules:

- High failure probability (> 0.8)
 - Low SoH ($< 75\%$)
 - Service downtime
 - Database connection failures
-

Deployment Considerations

Production Readiness

1. Kubernetes Migration:

- Convert Docker Compose to K8s manifests
- Use StatefulSets for databases
- Implement HorizontalPodAutoscaler

2. High Availability:

- TimescaleDB replication
- Kafka cluster (3+ nodes)
- Predictor service replicas

3. Security Hardening:

- TLS/SSL for all connections
- Secrets management (HashiCorp Vault)
- Network policies
- RBAC for services

4. CI/CD Pipeline:

- GitHub Actions / GitLab CI
- Automated testing

- Docker image building
- Staging environment

5. Backup Strategy:

- Daily database backups
 - Model artifact versioning
 - Configuration backups
 - Disaster recovery plan
-

Dataset Information

Source: <https://www.kaggle.com/datasets/datasetengineer/evoit-predictivemaint-dataset>

Expected Columns:

- SoC (State of Charge)
- SoH (State of Health)
- Battery_Voltage
- Battery_Current
- Battery_Temperature
- Charge_Cycles
- Power_Consumption
- RUL (Remaining Useful Life) - Target
- Failure_Probability - Target

Preprocessing:

- Handle missing values (forward fill for time series)
 - Remove outliers (IQR method)
 - Feature scaling (StandardScaler)
 - No data augmentation needed (large dataset)
-

Expected Output

After implementation, the system should:

1. Generate realistic battery telemetry every 2 seconds
2. Store time-series data in TimescaleDB
3. Stream data via Kafka and MQTT
4. Predict RUL with 90%+ accuracy ($R^2 > 0.90$)
5. Classify failure probability
6. Update predictions every 5 seconds
7. Expose Prometheus metrics
8. Display live Grafana dashboards
9. Track experiments in MLflow
10. Run containerized services

Visual Outputs:

- Real-time line charts of SoC, SoH, Temperature
- RUL gauge showing remaining cycles
- Failure risk percentage
- 3D battery visualization (bonus)
- Alert notifications

Additional Resources

Configuration Files to Create

Mosquitto Config (`infra/mosquitto/mosquitto.conf`):

```
listener 1883
allow_anonymous true
listener 9002
protocol websockets
```

Grafana Dashboard JSON (`infra/grafana/provisioning/dashboards/battery_dashboard.json`): Create panels for:

- Time series: SoC, SoH, Temperature, Voltage, Current
- Gauge: RUL, Failure Probability
- Stat: Charge Cycles, Power Consumption

Success Criteria

The implementation is complete when:

1. All 7 Docker services start successfully
 2. Models train with $R^2 > 0.90$ for RUL
 3. Simulator publishes to all 3 channels (DB, Kafka, MQTT)
 4. Predictor makes predictions every 5 seconds
 5. Grafana displays 8+ panels with live data
 6. Prometheus scrapes metrics from predictor
 7. MLflow tracks training experiments
 8. System runs continuously without crashes
 9. Documentation is complete
 10. Code passes linting (pylint, mypy)
-

Copilot-Specific Instructions

When implementing this project with GitHub Copilot:

1. **Start with infrastructure:** Generate docker-compose.yml first
2. **Database schema next:** Create SQL initialization script
3. **Training pipeline:** Implement model training with MLflow integration
4. **Simulator:** Build telemetry generator with realistic physics
5. **Predictor:** Create inference service with Prometheus metrics
6. **Configuration files:** Generate Grafana, Prometheus configs
7. **Documentation:** Auto-generate README with architecture diagrams
8. **Tests:** Create unit tests for all modules

Use these Copilot comments in your code:

```
python
```

```
# TODO: Implement EV battery telemetry simulator with charge/discharge cycles
# TODO: Train XGBoost model for RUL prediction with R2 > 0.90
# TODO: Create TimescaleDB connection with retry logic and connection pooling
# TODO: Publish telemetry to Kafka topic 'ev-telemetry' with error handling
# TODO: Expose Prometheus metrics on port 9100 for monitoring
```

Final Notes

This is a **production-grade** system designed for:

- Real-world EV battery monitoring
- Predictive maintenance applications
- IoT data pipeline demonstrations
- ML Ops best practices
- Time-series database optimization

The architecture is **scalable** and **maintainable**, following industry standards for observability, reliability, and performance.

Estimated Development Time: 2-3 days with Copilot assistance

License: MIT (or your choice) **Author:** Your Name **Last Updated:** October 20, 2025